

Mutation Operators for Concurrent Java (J2SE 5.0)¹

Jeremy S. Bradbury, James R. Cordy, Juergen Dingel
School of Computing, Queen's University
Kingston, Ontario, Canada
{bradbury, cordy, dingel}@cs.queensu.ca

Abstract

The current version of Java (J2SE 5.0) provides a high level of support for concurrency in comparison to previous versions. For example, programmers using J2SE 5.0 can now achieve synchronization between concurrent threads using explicit locks, semaphores, barriers, latches, or exchangers. Furthermore, built-in concurrent data structures such as hash maps and queues, built-in thread pools, and atomic variables are all at the programmer's disposal.

We are interested in using mutation analysis to evaluate, compare and improve quality assurance techniques for concurrent Java programs. Furthermore, we believe that the current set of method mutation operators and class operators proposed in the literature are insufficient to evaluate concurrent Java source code because the majority of operators do not directly mutate the portions of code responsible for synchronization. In this paper we will provide an overview of concurrency constructs in J2SE 5.0 and a new set of concurrent mutation operators. We will justify the operators by categorizing them with an existing bug pattern taxonomy for concurrency. Most of the bug patterns in the taxonomy have been used to classify real bugs in a benchmark of concurrent Java applications.

1 Introduction

As a result of advances in hardware technology (e.g. multi-core processors) a number of practitioners and researchers have advocated the need for concurrent software development [14]. Unfortunately, developing correct concurrent code is much more difficult than developing correct sequential code. The difficulty in programming concurrently is due to the many different, possibly unexpected, executions of the program. Reasoning about all possible interleavings in a program and ensuring that interleavings do

not contain bugs is non-trivial. Edward A. Lee discussed concurrency bugs in a recent paper [9]:

"I conjecture that most multithreaded-general purpose applications are so full of concurrency bugs that - as multicore architectures become commonplace - these bugs will begin to show up as system failures."

The presence of bugs in concurrent code can have serious consequences including deadlock, starvation, livelock, dormancy, and incoincidence (calls occurring at the wrong time) [11].

We are interested in using mutation to evaluate, compare, and improve quality assurance techniques for concurrent Java. The use of mutation with Java has been proposed in previous work – for instance the MuJava tool [13]. MuJava includes two general types of mutation operators for Java: method level operators [7, 13] and class level operators [12]. The method level operators include modifications to statements (e.g., statement deletion) and modifications to operands and operators in expressions (e.g., arithmetic operator insertion). The class level operators are related to inheritance (e.g., super keyword deletion), polymorphism (e.g., cast type change), and Java-specific features. **In general, the method and class level mutation operators do not directly mutate the synchronization portions of the source code in Java (J2SE 5.0) that handle concurrency.** Furthermore, we conjecture that additional operators are needed in order to provide a more comprehensive set of operators that can truly reflect the types of bugs that often occur in concurrent programs. In this paper we present a set of concurrent operators for Java (J2SE 5.0). We believe our new set of concurrency mutation operators used in conjunction with existing method and class level operators provide a more comprehensive set of mutation metrics for the comparison and improvement of quality assurance testing and analysis for concurrency.

In the next section (Section 2) we will provide an overview of the support for concurrency in Java (J2SE 5.0). In Section 3 we provide an overview of real concurrency

¹This work is supported by the Natural Sciences and Engineering Research Council of Canada (NSERC).

bug patterns which we will use to classify our concurrency mutation operators and demonstrate that the set of operators is both comprehensive and representative of real bugs. The set of mutation operators for concurrency and the bug pattern classification are presented in Section 4. Finally in Section 5 we provide our conclusions and an overview of our future work on using our new mutation operators.

2 Java Concurrency

Threads. Java concurrency is built around the notion of multi-threaded programs. The Java documentation defines a thread as “...a thread of execution in a program.”² A typical thread is created and then started using the `start()` method and will be terminated once it has finished running. While a thread is alive it can often alternate between being runnable and not runnable. A number of methods exist that can affect the status of a thread:

- `sleep()`: will cause the current thread to become not runnable for a certain amount of time.
- `yield()`: will cause the current thread that is running to pause (temporarily).
- `join()`: will cause the caller thread to wait for a target thread to terminate.
- `wait()`: will cause the caller thread to wait until a condition is satisfied. Another thread notifies the caller that a condition is satisfied using the `notify()` or `notifyAll()` method.

Synchronization. Prior to J2SE 5.0, Java provided support for concurrency primarily through the use of the synchronized keyword. Java supports both synchronization methods and synchronization blocks. Additionally, synchronization blocks can be used in combination with implicit monitor locks.

Other Concurrency Mechanisms. In J2SE 5.0, additional mechanisms to support concurrency were added as part of `java.util.concurrent`:

- **Explicit Lock:** Provides the same semantics as the implicit monitor locks but provides additional functionality such as timeouts during lock acquisition.
- **Semaphore:** Maintains a set of permits that restrict the number of threads accessing a resource. A Semaphore with one permit acts the same as a Lock.
- **Latch:** Allows threads from a set to wait until other threads complete a set of operations.
- **Barrier:** A point at which threads from a set wait until all other threads reach the point.
- **Exchanger:** Allows for the exchange of objects between two threads at a given synchronization point.

²`java.lang.Thread` documentation

Built-in Concurrent Data Structures. To reduce the overhead of developing concurrent data structures, J2SE 5.0 provides a number of collection types including `ConcurrentHashMap` and five different `BlockingQueues`.

Built-in Thread Pools. J2SE 5.0 provides a built-in `FixedThreadPool` and an unbounded `CachedThreadPool`.

Atomic Variables. The `java.util.concurrent.atomic` package includes a number of atomic variables that can be used in place of synchronization: `AtomicInteger`, `AtomicIntegerArray`, `AtomicLong`, `AtomicLongArray`, `AtomicBoolean`, `AtomicReference` and `AtomicReferenceArray`. Each atomic variable type contains new methods to support concurrency. For example, `AtomicInteger` contains methods such as `addAndGet()`, `getAndSet()` and others.

3 Bug Patterns for Java Concurrency

Farchi, Nir, and Ur have developed a bug pattern taxonomy for Java concurrency [6]. The bug patterns are based on common mistakes programmers make when developing concurrent code in practice. Furthermore, the taxonomy has been expanded and used to classify bugs in an existing public domain concurrency benchmark maintained by IBM Research [5]. The benchmark contains 40 programs ranging in size from 57 to 17000 loc. Programs in the benchmark are from a variety of sources including student created programs, tool developer programs, open source programs, and a commercial product. In our attempt to develop a comprehensive set of concurrency mutation operators we will later classify our operators with respect to the bug patterns taxonomy. Since this bug pattern taxonomy was developed prior to J2SE 5.0 we have had to add some additional patterns that occur in concurrency constructs not available at the time the taxonomy was proposed. We distinguish between the original bug patterns(*), the added bug patterns also used in the benchmark classification(**) and new patterns that we are including (+):

- **Nonatomic operations assumed to be atomic bug pattern.*** “...an operation that “looks” like one operation in one programmer model (e.g., the source code level of the programming language), but actually consists of several unprotected operations at the lower abstraction levels” [6]. In this paper we also include nonatomic floating point operations** in this pattern.
- **Two-state access bug pattern.*** “Sometimes a sequence of operations needs to be protected but the programmer wrongly assumes that separately protecting each operation is enough” [6].
- **Wrong lock or no lock bug pattern.*** “A code segment is protected by a lock but other threads do not obtain the same lock instance when executing. Either these other threads do not obtain a lock at all or they

obtain some lock other than the one used by the code segment” [6].

- **Double-checked lock bug pattern.*** “When an object is initialized, the thread local copy of the objects field is initialized but not all object fields are necessarily written to the heap. This might cause the object to be partially initialized while its reference is not null” [6].
- **The sleep() bug pattern.*** “The programmer assumes that a child thread should be faster than the parent thread in order that its results be available to the parent thread when it decides to advance. Therefore, the programmer sometimes adds an ‘appropriate’ sleep() to the parent thread. However, the parent thread may still be quicker in some environment. The correct solution would be for the parent thread to use the join() method to explicitly wait for the child thread” [6].
- **Losing a notify bug pattern.*** “If a notify() is executed before its corresponding wait(), the notify() has no effect and is “lost” ... the programmer implicitly assumes that the wait() operation will occur before any of the corresponding notify() operations” [6].
- **Other missing or nonexistent signals.**+ This pattern generalizes the losing a notify bug pattern to all other signals. For example, at a barrier the await() method has to be called by a set number of threads before the program can proceed. If an await() from one thread never occurs then all of threads at the barrier may be stuck waiting.
- **Notify instead of notify all bug pattern.**** If a notify() is executed instead of notifyAll() then threads with some of its corresponding wait() calls will not be notified [10].
- **A “blocking” critical section bug pattern.*** “A thread is assumed to eventually return control but it never does” [6].
- **The orphaned thread bug pattern.*** “If the master thread terminates abnormally, the remaining threads may continue to run, awaiting more input to the queue and causing the system to hang” [6].
- **The interference bug pattern.**** A pattern in which “...two or more concurrent threads access a shared variable and when at least one access is a write, and the threads use no explicit mechanism to prevent the access from being simultaneous.” [11]. The interference bug pattern can also be generalized from classic data race interference to include high level data races** which deal “...with accesses to sets of fields which are related and should be accessed atomically” [1].
- **The deadlock (deadly embrace) bug pattern.**** “...a situation where two or more processes are unable to proceed because each is waiting for one of the others to do something in a deadlock cycle ... For example,

this occurs when a thread holds a lock that another thread desires and vice-versa” [11].

- **Starvation bug pattern.**+ This bug occurs when their is a failure to “...allocate CPU time to a thread. This may be due to scheduling policies...” [8]. For example, an unfair lock acquisition scheme might cause a thread never to be scheduled.
- **Resource exhaustion bug pattern.**+ “A group of threads together hold all of a finite number of resources. One of them needs additional resources but no other thread gives one up” [8].
- **Incorrect count initialization bug pattern.**+ This pattern occurs when there is an incorrect initialization in a barrier for the number of parties that must be waiting for the barrier to trip, or an incorrect initialization of the number of threads required to complete some action in a latch, or an incorrect initialization of the number of permits in a semaphore.

4 Concurrent Mutation Operators

We propose five categories of mutation operators for concurrent Java: **modify parameters of concurrent methods**, **modify the occurrence of concurrency method calls** (removing, replacing and exchanging), **modify keywords** (addition and removal), **switch concurrent objects**, and **modify critical regions** (shift, expand, shrink and split). The relationship between these general operator categories and the concurrency mechanisms provided in J2SE 5.0 is presented in Table 1 – which demonstrates that the operators provide coverage over the J2SE 5.0 concurrency mechanisms.

A complete list of the operators we will be presenting in this section is provided in Table 2. The mutant operators are designed specifically to represent mistakes that programmers may make when implementing concurrency. Therefore, many of the operators are specific only to concurrency methods, objects and keywords. We have tried to use context and knowledge about Java concurrency to make the operators as specific as possible in order to make concurrency mutation analysis more feasible by reducing the total number of mutants produced.

Readers familiar with method and class level mutation operators will notice that some of our mutation operators are special cases of existing mutation operators while others are new operators that have not been previously proposed. Other related work from the concurrency bug detection community includes a set of 18 hand-created concurrency mutants [10] for a previous version of Java that did not contain many of the concurrency mechanisms available in J2SE 5.0. We have compared our comprehensive set of operators with this work and found that our operators in combination with the method and class level operators subsume the manual mutants used in the previous work.

Java (J2SE 5.0) Concurrency Mutation Operator Categories	Threads	Synchronization methods	Synchronization statements	Synchronization with implicit monitor locks	Explicit locks	Semaphores	Barriers	Latches	Exchangers	Built-in concurrent data structures (e.g. queues)	Built-in thread pools	Atomic variables (e.g. LongInteger)
Modify Parameters of Concurrent Methods	✓	—	✓	✓	✓	✓	✓	✓	—	—	—	—
Modify the Occurrence of Concurrency Method Calls	✓	—	—	—	✓	✓	✓	✓	—	—	—	✓
Modify Keyword	—	✓	✓	✓	✓	—	—	—	—	—	—	—
Switch Concurrent Objects	—	—	—	—	✓	✓	✓	✓	✓	✓	✓	—
Modify Concurrent Region	—	✓	✓	✓	✓	✓	—	—	—	—	—	—

Table 1. The relationship between new mutation operators for concurrency and the concurrency features provided by J2SE 5.0

Operator Category	Concurrency Mutation Operators for Java (J2SE 5.0)
Modify Parameters of Concurrent Methods	MXT - Modify Method-X Time (<i>wait()</i> , <i>sleep()</i> , <i>join()</i> , and <i>await()</i> method calls)
	MSP - Modify Synchronized Block Parameter
	ESP - Exchange Synchronized Block Parameters
	MSF - Modify Semaphore Fairness
	MXC - Modify Permit Count in Semaphore and Modify Thread Count in Latches and Barriers
	MBR - Modify Barrier Runnable Parameter
Modify the Occurrence of Concurrency Method Calls	RTXC - Remove Thread Method-X Call (<i>wait()</i> , <i>join()</i> , <i>sleep()</i> , <i>yield()</i> , <i>notify()</i> , <i>notifyAll()</i> Methods)
	RCXC - Remove Concurrency Mechanism Method-X Call (<i>methods in Locks, Semaphores, Latches, Barriers, etc.</i>)
	RNA - Replace NotifyAll() with Notify()
	RJS - Replace Join() with Sleep()
	ELPA - Exchange Lock/Permit Acquisition
	EAN - Exchange Atomic Call with Non-Atomic
Modify Keyword	ASTK - Add Static Keyword to Method
	RSTK - Remove Static Keyword from Method
	RSK - Remove Synchronized Keyword from Method
	RSB - Remove Synchronized Block
	RVK - Remove Volatile Keyword
	RFU - Remove Finally Around Unlock
Switch Concurrent Objects	RXO - Replace One Concurrency Mechanism-X with Another (<i>Locks, Semaphores, etc.</i>)
	EELO - Exchange Explicit Lock Objects
Modify Critical Region	SHCR - Shift Critical Region
	SKCR - Shrink Critical Region
	EXCR - Expand Critical Region
	SPCR - Split Critical Region

Table 2. Concurrency mutation operators for Java

4.1 Modify parameters of concurrent method

These operators involve modifying the parameters of methods for thread and concurrency classes. Some of the

method level mutation operators that modify operands are similar to the operators proposed here.

4.1.1 MXT - Modify Method-X Timeout

The MXT operator can be applied to the *wait()*, *sleep()*, and *join()* method calls (introduced in Section 2) that include an optional timeout parameter. For example, in Java a call to *wait()* with the optional timeout parameter will cause a thread to no longer be runnable until a condition is satisfied or a timeout has occurred. The MXT replaces the timeout parameter, *t*, of the *wait()* method by some appropriately chosen fraction or multiple of *t* (e.g., *t/2* and *t * 2*). We could replace the timeout parameter by a variable of an equivalent type however since we know that the parameter represents a time value it is just as meaningful to mutate the method to both increase and decrease the time by a factor of 2.

Original Code:

```
long time = 10000;
try {
    wait (time);
} catch ...
```

MXT Mutant:

```
long time = 10000;
try {
    wait (time * 2);
    // or time / 2
} catch ...
```

The MXT operator with the *wait()* method is most likely to result in an interference bug or a data race. The MXT operator with the *sleep()* and *join()* methods is most likely to result in the *sleep()* bug pattern. For example, in a situation where a *sleep()* or *join()* is used by a caller thread to wait for another thread, reducing the time may cause the caller thread to not wait long enough for the other thread to complete.

The MXT operator can also be applied to the optional timeout parameter in *await()* method calls. Both barriers and latches have an *await()* method. In barriers the *await()*

method is used to cause a thread to wait until all threads have reached the barrier. In latches the `await()` method is used by threads to wait until the latch has finished counting down, that is until all operations in a set are complete. The MXT operator when applied to an `await()` method call will most likely result in an interference bug.

4.1.2 MSP - Modify Synchronized Block Parameter

Common parameters for a synchronized block include the `this` keyword, indicating that synchronization occurs with respect to the instance object of the class, and implicit monitor objects. If the keyword `this` or an object is used as a parameter for a synchronized block we can replace the parameter by another object or the keyword `this`. For example:

Original Code:

```
private Object lock1 = new Object();
private Object lock2 = new Object();
...
public void methodA() {
    synchronized(lock1) { ... }
}
...
```

MSP Mutant:

```
private Object lock1 = new Object();
private Object lock2 = new Object();
...
public void methodA() {
    synchronized(lock2) { ... }
}
...
```

Another MSP Mutant:

```
private Object lock1 = new Object();
private Object lock2 = new Object();
...
public void methodA() {
    synchronized(this) { ... }
}
...
```

The MSP operator will result in the wrong lock bug pattern.

4.1.3 ESP - Exchange Synchronized Block Parameters

If a critical region is guarded by multiple synchronized blocks with implicit monitor locks the ESP operator exchanges two adjacent lock objects. For example:

Original Code:

```
private Object lock1 = new Object();
private Object lock2 = new Object();
...
public void methodA() {
    synchronized(lock1) {
        synchronized(lock2) { ... }
    }
}
...
public void methodB() {
    synchronized(lock1) {
        synchronized(lock2) { ... }
    }
}
...
```

ESP Mutant:

```
private Object lock1 = new Object();
private Object lock2 = new Object();
...
public void methodA() {
    synchronized(lock2) {
        synchronized(lock1) { ... }
    }
}
...
public void methodB() {
    synchronized(lock1) {
        synchronized(lock2) { ... }
    }
}
...
```

The ESP mutation operator can result in a wrong lock bug because exchanging two adjacent locks will cause the locks to be acquired at incorrect times for incorrect critical regions. The ESP operator can also cause a classic deadlock (via deadly embrace) bug to occur as is the case in the above example.

4.1.4 MSF - Modify Semaphore Fairness

Recall in Section 2 that a semaphore maintains a set of permits for accessing a resource. In the constructor of a Semaphore there is an optional parameter for a boolean fairness setting. When the fairness setting is not used the default fairness value is `false` which allows for unfair permit acquisition. If the fairness parameter is a constant then the MSF operator is a special case of the Constant Replacement (CRP) method level operator and replaces a `true` value with `false` and a `false` value with `true`. In the case that a boolean variable is used as a parameter we simply negate it.

A potential consequence of expecting a semaphore to be fair when in fact it is not is that there is a potential for starvation because no guarantees about permit acquisition ordering can be given. In fact, when a semaphore is unfair any thread that invokes the Semaphore's `acquire()` method to obtain a permit may receive one prior to an already waiting thread - this is known as *barging*³.

Original Code:

```
int permits = 10;
private final Semaphore sem
    = new Semaphore(permits, true);
...
```

MSF Mutant:

```
int permits = 10;
private final Semaphore sem
    = new Semaphore(permits, false);
...
```

³java.util.concurrent documentation

4.1.5 MXC - Modify Concurrency Mechanism-X Count

The MXC operator is applied to parameters in three of Java's concurrency mechanisms: Semaphores, Latches, and Barriers. A latch allows a set of threads to countdown a set of operations and a barrier allows a set of threads to wait at a point until a number of threads reach that point. The count being modified in Semaphores is the set of permits, and in Latches and Barriers it is the number of threads. We will next provide an example of the MXC operator for Semaphores. For examples involving Latches and Barriers see our technical report [3].

The constructor of the Semaphore class has a parameter that refers to the maximum number of available permits that are used to limit the number of the threads accessing the shared resource. Access is acquired using the `acquire()` method and released using the `release()` method. Both the `acquire()` and `release()` method calls have optional count parameters referring to the number of permits being acquired or released. The MXC operator modifies the number of permits, p , in calls to these methods by decrementing ($p--$) and incrementing ($p++$) it by 1. For example:

Original Code:

```
int permits = 10;
private final Semaphore sem
    = new Semaphore (permits, true);
...
```

MSC Mutant:

```
int permits = 10;
private final Semaphore sem
    = new Semaphore (permits --, true);
...
```

A potential bug that can occur from modifying permit counts in Semaphores or number of threads in Latches and Barriers is resource exhaustion. In the above example if the total number of permits had been one then decrementing the number of permits by 1 would have lead to a situation where no permits were ever available. Another bug could occur if we increased the number of permits acquired by the `acquire()` method but did not increase the count in the `release()` method which could eventually exhaust the resources. In this case we could end up with a blocking critical section bug once all of the permits were held but not released.

4.1.6 MBR - Modify Barrier Runnable Parameter

The `CyclicBarrier` constructor has a parameter that is an optional runnable thread that can happen after all the threads complete and reach the barrier. The MBR operator modifies the runnable thread parameter by removing it if it is present. This is a special case of the method level mutation operator, statement deletion (SDL). For example:

Original Code:

```
int i=10;
CyclicBarrier barrier1
    = new CyclicBarrier(i,
        new Runnable() {
            public void run() {
            }
        });
...
```

MBR Mutant:

```
int i=10;
CyclicBarrier barrier1
    = new CyclicBarrier(i);
...
```

An example of a bug caused by the MBR operator is missed or nonexistent signals if some signal calls were present in the runnable thread.

4.2 Modify the occurrence of concurrency method calls: remove, replace, and exchange

This class of operators is primarily interested in modifying calls to thread methods and methods of concurrency mechanism classes. Examples of modifications include removal of a method call and replacement or exchange of a method call with a different but similar method call. The operators that remove method calls are special cases of the method level operator: Statement Deletion (SDL).

4.2.1 RTXC - Remove Thread Method-X Call

The RTXC operator removes calls to the following methods: `wait()`, `join()`, `sleep()`, `yield()`, `notify()`, and `notifyAll()`. Removing the `wait()` method can cause potential interference, removing the `join()` and `sleep()` methods can cause the `sleep()` bug pattern, and removing the `notify()` and `notifyAll()` method calls is an example of losing a notify bug. We will now provide an example of the RTXC operator used to remove a `wait()` method call.

Original Code:

```
try {
    wait ();
} catch ...
```

RTXC Mutant:

```
try {
    //removed
    //wait();
} catch ...
```

4.2.2 RCXC - Remove Concurrency Mechanism Method-X Call

The RCXC operator can be applied to the following concurrency mechanisms: Locks (`lock()`, `unlock()`), Condition (`signal()`, `signalAll()`), Semaphore (`acquire()`, `release()`), Latch (`countDown()`), and `ExecutorService` (e.g., `submit()`). For details on each method as well as the application of the RCXC operator to each method see [3]. In this paper we will only discuss the RCXC operator when using locks. In

a ReentrantLock or a ReentrantReadWriteLock a call to the unlock() method attempts to release the lock. The RCXC operator removes this call thus the lock is not released. This is an example of a blocking critical section bug. For example:

Original Code: <pre>private Lock lock1 = new ReentrantLock(); ... lock1.lock(); try { ... } finally { lock1.unlock(); } ...</pre>	RCXC Mutant: <pre>private Lock lock1 = new ReentrantLock(); ... lock1.lock(); try { ... } finally { //removed lock1.unlock(); } ...</pre>
---	---

4.2.3 RNA - Replace NotifyAll() with Notify() RJS - Replace Join() with Sleep()

The RNA operator replaces a notifyAll() with a notify() and is an example of the notify instead of notify all bug pattern.

Original Code: <pre>... notifyAll();...</pre>	RNA Mutant: <pre>... notify();...</pre>
---	---

The RJS operator replaces a join() with a sleep() and is an example of the sleep() bug pattern.

Original Code: <pre>... join();...</pre>	RJS Mutant: <pre>... sleep(10000);...</pre>
--	---

4.2.4 ELPA - Exchange Lock/Permit Acquisition

In a Semaphore the acquire(), acquireUninterruptibly() and tryAcquire() methods can be used to obtain one or more permits to access a shared resource. The ELPA operator exchanges one method for another which can lead to potential timing changes as well as starvation. For example, an acquire() method will try and obtain one or more permits and will block and wait until the permit or permits become available. If the thread that invoked the acquire() method is interrupted it will no longer continue to block and wait. If the acquire() method invocation is changed to acquireUninterruptibly() it will behave exactly the same except it can no longer be interrupted. Thus in situations where the semaphore is unfair or if for other reasons the number of requested permits never becomes available the thread that invoked the acquireUninterruptibly() will stay dormant and wait. If an acquire() method invocation is changed to a tryAcquire() then a permit will be acquired if one is available otherwise the thread will not block and wait. tryAcquire() will acquire a permit or permits unfairly even if the fairness setting is set to fair. Use of tryAcquire() may cause starvation for threads waiting for permits.

Original Code:

```
int permits = 10;
private final Semaphore sem
    = new Semaphore (permits, true);
...
sem.acquire();
...
```

ELPA Mutant:

```
int permits = 10;
private final Semaphore sem
    = new Semaphore (permits, true);
...
sem.acquireUninterruptibly();
...
```

Another ELPA Mutant:

```
int permits = 10;
private final Semaphore sem
    = new Semaphore (permits, true);
...
sem.tryAcquire();
...
```

The ELPA operator can also be applied to the lock(), lockInterruptibly(), tryLock() method calls with Locks.

4.2.5 SAN - Switch Atomic Call with Non-Atomic

A call to the getAndSet() method in an atomic variable class is replaced by a call to the get() method and a call to the set() method. The effect of this replacement is that the combined get and set commands are no longer atomic. For example:

Original Code: <pre>AtomicInteger int1 = 15; ... int oldVal = int1.getandSet(40); ...</pre>	SAN Mutant: <pre>AtomicInteger int1 = 15; ... int oldVal = int1.get(); int1.set(40); ...</pre>
---	--

4.3 Modify keywords: add and remove

We consider what happens when we add and remove keywords such as static, synchronized, volatile, and finally.

4.3.1 ASTK - Add Static Keyword to Method RSTK - Remove Static Keyword from Method

The static keyword used for a synchronized method indicates that the method is synchronized using the class object not the instance object. The ASTK operator adds static to non-static synchronized methods and the RSTK removes static from static synchronized methods. Since the addition or removal of the static keyword causes synchronization to occur on the class or instance object the ASTK and RSTK operators are both examples of the wrong lock bug pattern.

Original Code:

```
public synchronized void aMethod() { ... }
```

ASTK Mutant:

```
public static synchronized void aMethod() { ... }
```

Original Code:

```
public static synchronized void bMethod() { ... }
```

RSTK Mutant:

```
public synchronized void bMethod() { ... }
```

4.3.2 RSK - Remove Synchronized Keyword from Method

The synchronized keyword is important in defining concurrent methods and the omission of this keyword is a plausible mistake that a programmer might make when writing concurrent source code. The RSK operator removes the synchronized keyword from a synchronized method and causes a potential no lock bug. For example:

Original Code:

```
public synchronized void aMethod() { ... }
```

RSK Mutant:

```
public void aMethod() { ... }
```

4.3.3 RSB - Remove Synchronized Block

Similar to the RSK operator, the RSB operator removes the synchronized keyword from around a statement block which can cause a no lock bug.

Original Code:

```
synchronized (this) {  
    <statement_c1>  
}
```

RSB Mutant:

```
{  
    <statement_c1>  
}
```

4.3.4 RVK - Remove Volatile Keyword

The volatile keyword is used with a shared variable and prevents operations on the variable from being reordered in memory with other operations. In the below example we remove the volatile keyword from a shared long variable. If a long variable, which is 64-bit, is not declared volatile then reads and writes will be treated as two 32-bit operations instead of one operation. Therefore, the RVK operator can cause a situation where a nonatomic operation is assumed to be atomic.

Original Code:

```
volatile long x;
```

RVK Mutant:

```
long x;
```

4.3.5 RFU - Remove Finally Around Unlock

The finally keyword is important in releasing explicit locks. In the below example, finally ensures that the unlock() method call will occur after a try block regardless of whether or not an exception is thrown. If finally is removed the unlock() will not occur in the presence of an exception and cause a blocking critical section bug.

Original Code:

```
private Lock lock1  
    = new ReentrantLock();  
...  
lock1.lock();  
try {  
    ...  
} finally {  
    lock1.unlock();  
}  
...
```

RFU Mutant:

```
private Lock lock1  
    = new ReentrantLock();  
...  
lock1.lock();  
try {  
    ...  
} lock1.unlock();  
...
```

4.4 Switch concurrent objects

When multiple instances of the same concurrent class type exist we can replace one concurrent object with the other.

4.4.1 RXO - Replace One Concurrency Mechanism-X with Another

When two instances of the same concurrency mechanism exist we replace a call to one with the other. Due to space we will only consider the replacement of Locks:

Original Code:

```
private Lock lock1  
    = new ReentrantLock();  
private Lock lock2  
    = new ReentrantLock();  
...  
lock1.lock();  
...
```

RXO Mutant:

```
private Lock lock1  
    = new ReentrantLock();  
private Lock lock2  
    = new ReentrantLock();  
...  
lock2.lock();  
...
```

We can also apply the RXO operator when 2 or more objects exist of type Semaphore, CountdownLatch, CyclicBarrier, Exchanger, and more. For details on the RXO operator with these other mechanisms see [3].

4.4.2 EELO - Exchange Explicit Lock Object

We have already seen the exchanging of two implicit lock objects in a synchronized block and the potential for deadlock (Section 4.1.3). The EELO operator is identical only it exchanges two explicit lock object instances:

Original Code:

```
private Lock lock1  
    = new ReentrantLock();  
private Lock lock2  
    = new ReentrantLock();  
...  
lock1.lock();  
...  
lock2.lock();  
...  
finally {  
    lock2.unlock();  
}  
...  
finally {  
    lock1.unlock();  
}  
...
```

EELO Mutant:

```
private Lock lock1  
    = new ReentrantLock();  
private Lock lock2  
    = new ReentrantLock();  
...  
lock2.lock();  
...  
lock1.lock();  
...  
finally {  
    lock2.unlock();  
}  
...  
finally {  
    lock1.unlock();  
}  
...
```


4.5 Modify critical region : shift, expand, shrink and split

The modify critical region operators cause the modification of the critical region by moving statements both inside and outside the region and by dividing the region into multiple regions.

4.5.1 SHCR - Shift Critical Region

Shifting a critical region up or down can potentially cause interference bugs by no longer synchronizing access to a shared variable. An example of shifting a synchronized block up is provided below. The SHCR operator can also be applied to shift up or down critical regions using other concurrency mechanisms. We also provide an example in the extended technical report version of this paper of shifting the critical region using explicit locks [3].

Original Code:

```
<statement n1>
<statement n2>
synchronized (this) {
    //critical region
    <statement c1>
    <statement c2>
}
<statement n3>
<statement n4>
...
```

SHCR Mutant:

```
<statement n1>
<statement n2>
//critical region
<statement c1>
synchronized (this) {
    <statement c2>
    <statement n3>
}
<statement n4>
...
```

4.5.2 EXCR - Expand Critical Region

Expanding a critical region to include statements above and below the statements required to be in the critical region can cause performance issues by unnecessarily reducing the degree of concurrency. For example:

Original Code:

```
<statement n1>
<statement n2>
synchronized (this) {
    //critical region
    <statement c1>
    <statement c2>
}
<statement n3>
<statement n4>
...
```

EXCR Mutant:

```
<statement n1>
synchronized (this) {
    <statement n2>
    //critical region
    <statement c1>
    <statement c2>
    <statement n3>
}
<statement n4>
...
```

The EXCR operator can also cause correctness issues and consequences such as deadlock when an expanded critical region overlaps with or subsumes another critical region.

4.5.3 SKCR - Shrink Critical Region

Shrinking a critical region will have similar consequences (interference) to shifting a region since both the SHCR and

SKCR operators move statements that require synchronization outside the critical section. Below we provide an example of the SKCR operator using a Lock.

Original Code:

```
private Lock lock1
= new ReentrantLock();
...
public void m1 () {
<statement n1>
lock1.lock();
try {
    //critical region
    <statement c1>
    <statement c2>
    <statement c3>
} finally {
    lock1.unlock();
}
<statement n2>
...
```

SKCR Mutant:

```
private Lock lock1
= new ReentrantLock();
...
public void m1 () {
<statement n1>
//critical region
<statement c1>
lock1.lock();
try {
    <statement c2>
} finally {
    lock1.unlock();
}
<statement c3>
<statement n2>
...
```

4.5.4 SPCR - Split Critical Region

Unlike the SHCR or SKCR operators, splitting a critical region into two regions will not cause statements to move outside of the critical region. However, the consequences of splitting a critical region into 2 regions is potentially just as serious since a split may cause a set of statements that were meant to be atomic to be nonatomic. For example, in between the two split critical regions another thread might be able to acquire the lock for the region and modify the value of shared variables before the second half of the old critical region is executed.

Original Code:

```
<statement n1>
synchronized (this) {
    //critical region
    <statement c1>
    <statement c2>
}
<statement n2>
...
```

SPCR Mutant:

```
<statement n1>
synchronized (this) {
    //critical region
    <statement c1>
}
synchronized (this) {
    <statement c2>
}
<statement n2>
...
```

4.6 Summary

In the above subsections we have provided an overview of concurrency mutation operators for Java (J2SE 5.0). For more details on the operators as well as the relationship between our new operators and the existing method and class level operators please see our extended technical report version of this paper [3]. In our discussion of each operator we have briefly mentioned the bug pattern that relates to that operator. Table 3 provides a summary of this relationship and shows that the operators we propose are examples of real bug patterns. Overall almost all of the bug patterns are covered by the operators demonstrating that the proposed concurrency operators are not only representative but

provide good coverage. The bug patterns that do not have mutation operators are typically more specific complex patterns and the development of general operators related to these patterns is not feasible.

Concurrency Bug Pattern	Mutation Operators
Nonatomic operations assumed to be atomic bug pattern	RVK, EAN
Two-stage access bug pattern	SPCR
Wrong lock or no lock bug pattern	MSP, ESP, EELO, SHCR, SKCR, EXCR, RSB, RSK, ASTK, RSTK, RCXC, RXO
Double-checked locking bug pattern	—
The sleep() bug pattern	MXT, RJS, RTXC
Losing a notify bug pattern	RTXC, RCXC
Notify instead of notify all bug pattern	RNA
Other missing or nonexistent signals bug pattern	MXC, MBR, RCXC
A “blocking” critical section bug pattern	RFU, RCXC
The orphaned thread bug pattern	—
The interference bug pattern	MXT, RTXC, RCXC
The deadlock (deadly embrace) bug pattern	ESP, EXCR, EELO, RXO
Starvation bug pattern	MSF, ELPA
Resource exhaustion bug pattern	MXC
Incorrect count initialization bug pattern	MXC

Table 3. Concurrency bug patterns vs. concurrency mutation operators

5 Conclusion

We have presented a set of concurrency mutation operators to be used as a metric in the comparison of different test suites and testing strategies for concurrent Java as well as different quality assurance tools for concurrency. Although we are primarily interested in concurrent mutation operators as comparative metrics we believe that these operators can also serve a role similar to method and class level mutation operators as both comparative metrics and coverage criteria. Our new concurrency operators should be viewed as a complement not a replacement for the existing operators used in tools like MuJava. For example, using the concurrency operators can cause direct concurrency bugs while using the method and class level operators can cause indirect concurrency bugs.

We believe that our concurrency operators are comprehensive and representative of real bugs. We have justified the operators by comparing them to a set of bug patterns that have been used to identify real bugs in concurrent Java programs. Additionally, our classification of concurrency operators shows that the operators are well distributed across the majority of bug patterns.

Currently we are implementing our concurrency muta-

tion operators in a source transformation language TXL [4]. Upon completion of our implementation we plan to validate the operators with our mutation analysis framework ExMAN [2]. We are interested in using our concurrency operators with the programs in the IBM concurrency benchmark to compare concurrency testing and model checking.

References

- [1] C. Artho, K. Havelund, and A. Biere. High-level data races. *Software Testing, Verification and Reliability (Special Issue VVEIS 2003)*, 13(4):207–227, 2003.
- [2] J. S. Bradbury, J. R. Cordy, and J. Dingel. ExMAN: A generic and customizable framework for experimental mutation analysis. In *Proc. of the Work. on Mutation Analysis (Mutation 2006)*, 2006.
- [3] J. S. Bradbury, J. R. Cordy, and J. Dingel. Mutation operators for concurrent Java J2SE 5.0. Technical Report 2006-520, Queen’s University, 2006.
- [4] J. Cordy, T. Dean, A. Malton, and K. Schneider. Source transformation in software engineering using the TXL transformation system. *J. of Information and Software Technology*, 44(13):827–837, 2002.
- [5] Y. Eytani and S. Ur. Compiling a benchmark of documented multi-threaded bugs. In *Proc. of PADTAD 2004*, pages 266–273, 2004.
- [6] E. Farchi, Y. Nir, and S. Ur. Concurrent bug patterns and how to test them. In *Proc. of IPDPS 2003*, 2003.
- [7] K. N. King and A. J. Offutt. A Fortran language system for mutation-based software testing. *Softw. Pract. Exper.*, 21(7):685–718, 1991.
- [8] D. Lea. *Concurrent Programming in JavaTM Second Edition*. Addison Wesley, 2000.
- [9] E. Lee. The problem with threads. *Computer*, 39(5):33–42, May 2006.
- [10] B. Long, R. Duke, D. Goldson, P. A. Strooper, and L. Wildman. Mutation-based exploration of a method for verifying concurrent Java components. In *Proc. of IPDPS 2004*, page 265, Apr. 2004.
- [11] B. Long, P. Strooper, and L. Wildman. A method for verifying concurrent Java components based on an analysis of concurrency failures. In *Concurrency Computat.: Pract. Exper. (in press)*, 2006.
- [12] Y.-S. Ma, Y.-R. Kwon, and J. Offutt. Inter-class mutation operators for Java. In *Proc. of ISSRE 2002*, pages 352–363. IEEE Computer Society Press, Nov. 2002.
- [13] Y.-S. Ma, J. Offutt, and Y.-R. Kwon. MuJava : An automated class mutation system. *J. of Software Testing, Verification and Reliability*, 15(2):97–133, Jun. 2005.
- [14] H. Sutter and J. Larus. Software and the concurrency revolution. *Queue*, 3(7):54–62, 2005.